maze-dataset

Michael Igorevich Ivanitskiy [1][¶]**, Aaron Sandoval** [1]**, Alex F. Spies** [2]**,
Tilman Räuker** [1]**, Brandon Knutson** [1]**, Cecilia Diniz Behn** [1]**, and Samy
Wu Fung** [1]

1 Colorado School of Mines, Department of Applied Mathematics and Statistics **2** Imperial College
London ¶ Corresponding author

Summary

Solving mazes is a classic problem in computer science and artificial intelligence, and humans
have been constructing mazes for thousands of years. Although finding the shortest path
through a maze is a solved problem, this very fact makes it an excellent testbed for studying how
machine learning algorithms solve problems and represent spatial information. We introduce
`maze-dataset`, a Python library for generating, processing, and visualizing datasets of mazes.
This library supports a variety of maze generation algorithms providing mazes with or without
loops, fully-connected or not, and many other variations. These generation algorithms can be
configured with various parameters, and the resulting mazes can be filtered to satisfy desired
properties. Also provided are tools for converting mazes to and from various formats suitable for
a variety of neural network architectures, such as rasterized images, tokenized text sequences,
and various visualizations. As well as providing a simple interface for generating, storing, and
loading these datasets, `maze-dataset` is extensively tested, type hinted, benchmarked, and
documented.
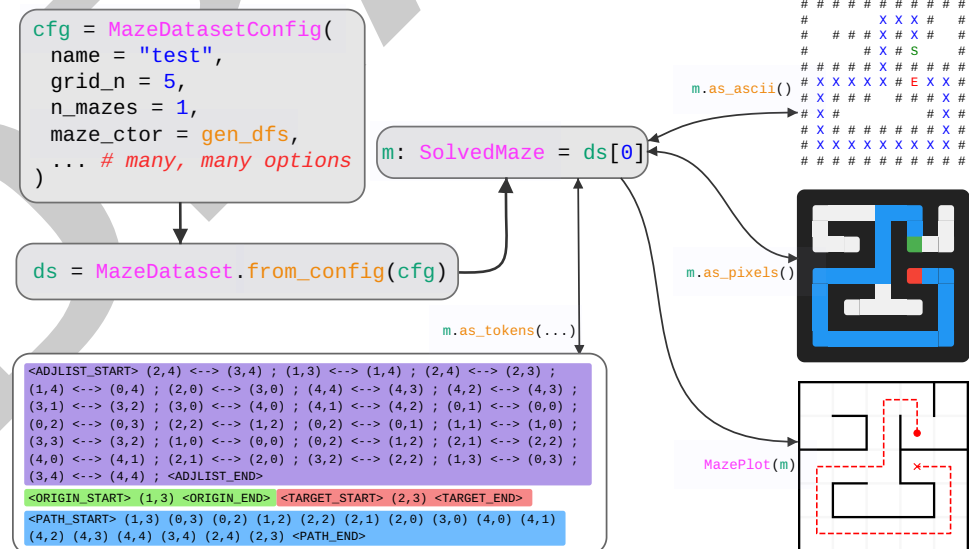
Figure 1: Usage of maze-dataset. We create a `MazeDataset` from a `MazeDatasetConfig`. This contains
`SolvedMaze` objects which can be converted to and from a variety of formats. Code in the image contains
clickable links to documentation. A variety of generated examples can be viewed here.

Statement of Need

The generation of mazes with a given algorithm is not inherently a complex task, but the ability to seamlessly switch out algorithms, modify algorithm parameters, or filter by desired properties all while preserving the ability to convert between different representations of the maze is not trivial. This library aims to greatly streamline the process of generating and working with datasets of mazes that can be described as subgraphs of an $n \times n$ lattice with boolean connections and, optionally, start and end points that are nodes in the graph. Furthermore, we place emphasis on a wide variety of possible text output formats aimed at evaluating the spatial reasoning capabilities of Large Language Models and other text-based transformer models.

For interpretability and behavioral research, algorithmic tasks offer benefits by allowing systematic data generation and task decomposition, as well as simplifying the process of circuit discovery (Räuker et al., 2023). Although mazes are well suited for these investigations, we have found that existing maze generation packages (Cobbe et al., 2019; Ehsan, 2022; Harries et al., n.d.; Németh, 2019; Schwarzschild, Borgnia, Gupta, Bansal, et al., 2021) do not support flexible maze generation algorithms that provide fine-grained control of generation parameters and the ability to easily transform between multiple representations of the mazes (Raster, Textual, Tokenized) for training and testing models.

Related Works

A multitude of public and open-source software packages exist for generating mazes (Ehsan, 2022; Németh, 2019; Schwarzschild, Borgnia, Gupta, Bansal, et al., 2021). However, nearly all of these packages generate and store mazes in a form that is not optimized for storage space or, more importantly, computer readability. The mazes produces by other packages are usually rasterized or in some form of image, rather than the underlying graph structure, and this makes it difficult to work with these datasets.

- Most prior works provide mazes in some kind of image or raster format, and we provide a variety of similar output formats:

    - `RasterizedMazeDataset`, utilizing `as_pixels()`, which can exactly mimic the outputs provided in easy-to-hard-data(Schwarzschild, Borgnia, Gupta, Bansal, et al., 2021) and can be configured to be similar to the outputs of Németh (2019)
    - `as_ascii()` provides a format similar to that used in (Oppenheim, 2018; Singla, 2023)
    - `MazePlot` provides a feature-rich plotting utility with support for multiple paths, heatmaps over positions, and more. This is similar to the outputs of (Alance AB, 2019; Ehsan, 2022; Guo et al., 2011; Nag, 2020)

- The text format provided by `SolvedMaze(...).as_tokens()` is similar to that of (Liu & Wu, 2023), but provides over 5.8 million unique formats for converting mazes to a text stream, detailed in subsection .

- For rigorous investigations of the response of a model to various distributional shifts, preserving metadata about the generation algorithm with the dataset itself is essential. To this end, our package efficiently stores the dataset along with its metadata in a single human-readable file (M. Ivanitskiy, n.d.). As far as we are aware, no existing packages do this reliably.

- Storing mazes as images is not only difficult to work with, but also inefficient. We use an extremely efficient method detailed in section .

- Our package is easily installable with source code freely available. It is extensively tested, type hinted, benchmarked, and documented. Many other maze generation packages lack this level of rigor and scope, and some (Ayaz et al., 2008) appear to simply no longer be accessible.

Features

Generation and Usage

Our package can be installed from PyPi via `pip install maze-dataset`, or directly from the git repository (Michael I. Ivanitskiy et al., 2023a).

To create a dataset, we first create a `MazeDatasetConfig` configuration object, which specifies the seed, number, and size of mazes, as well as the generation algorithm and its corresponding parameters. This object is passed to a `MazeDataset` class to create a dataset. Crucially, this `MazeDataset` mimics the interface of a PyTorch (Paszke et al., 2019) `Dataset`, and can thus be easily incorporated into existing data pre-processing and training pipelines, e.g., through the use of a `DataLoader` class.

```python
from maze_dataset import MazeDataset, MazeDatasetConfig, LatticeMazeGenerators
cfg: MazeDatasetConfig = MazeDatasetConfig(
    name="example",
    grid_n=3,
    n_mazes=32,
    maze_ctor=LatticeMazeGenerators.gen_dfs,
)
dataset: MazeDataset = MazeDataset.from_config(cfg)
```

When initializing mazes, further configuration options can be specified through the `from_config()` factory method as necessary. Options allow for saving/loading existing datasets instead of regenerating, and parallelization options for generation. Available maze generation algorithms are static methods of the `LatticeMazeGenerators` class and include generation algorithms based on randomized depth-first search, Wilson's algorithm (Wilson, 1996), percolation (Duminil-Copin, 2017; Fisher & Essam, 2004), Kruskal's algorithm (Kruskal, 1956), and others.

Furthermore, a dataset of mazes can be filtered to satisfy certain properties. Custom filters can be specified, and some filters are included in `MazeDatasetFilters`. For example, we can require a minimum path length from origin to target of 3:
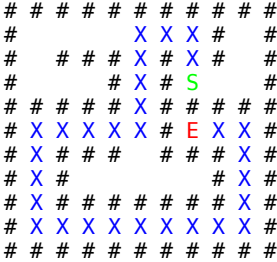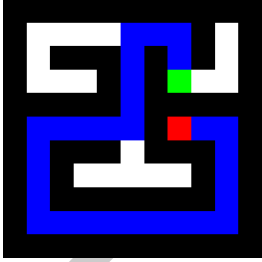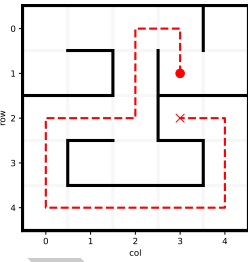
```python
dataset_filtered: MazeDataset = dataset.filter_by.path_length(min_length=3)
```

All implemented maze generation algorithms are stochastic by nature. For reproducibility, the seed parameter of `MazeDatasetConfig` may be set. In practice, we do not find that exact duplicates of mazes are generated with any meaningful frequency, even when generating large datasets.

Visual Output Formats

Internally, mazes are `SolvedMaze` objects, which have path information, and a connection list optimized for storing sub-graphs of a lattice. These objects can be converted to and from several formats to maximize their utility in different contexts.

`as_ascii()`	`as_pixels()`	`MazePlot()`
Simple text format for displaying mazes, useful for debugging in a terminal environment.	numpy array of `dtype=uint8` and shape (`height`, `width`, 3). The last dimension is RGB color.	feature-rich plotting utility with support for multiple paths, heatmaps over positions, and more.

Figure 2: Various output formats. Top row (left to right): ASCII diagram, rasterized pixel grid, and advanced display.

In previous work, maze tasks have been used with Recurrent Convolutional Neural Network (RCNN) derived architectures (Schwarzschild, Borgnia, Gupta, Huang, et al., 2021). To facilitate the use of our package in this context, we replicate the format of (Schwarzschild, Borgnia, Gupta, Bansal, et al., 2021) and provide the `RasterizedMazeDataset` class which returns rasterized pairs of (input, target) mazes as shown in Figure 3 below.

Figure 3: Input is the rasterized maze without the path marked (left), and provide as a target the maze with all but the correct path removed. Configuration options exist to adjust whether endpoints are included and if empty cells should be filled in.

Tokenized Output Formats

Autoregressive transformer models can be quite sensitive to the exact format of the input data, and may use delimiter tokens to perform reasoning steps (Pfau et al., 2024). To facilitate systematic investigation of this, we provide a variety of tokenized output formats.

We convert mazes to token sequences in two steps. First, the maze is stringified using `as_tokens()`. The `MazeTokenizerModular` class provides a powerful interface for configuring maze stringification behavior. Second, the sequence of strings is tokenized into integers using `encode()`. Tokenization uses a fixed vocabulary for simplicity. Mazes up to 50×50 are supported using unique tokens, and up to 128×128 when using coordinate tuple tokens.

There are many algorithms by which one might tokenize a 2D maze into a 1D format usable by autoregressive text models. Training multiple models on the encodings output from each of these algorithms may produce very different internal representations, learned solution algorithms, and levels of performance. To allow exploration of how different maze tokenization algorithms affect these models, the `MazeTokenizerModular` class contains a rich set of options

to customize how mazes are stringified. This class contains 19 discrete parameters, resulting in over 5.8 million unique tokenizers. But wait, there's more! There are 6 additional parameters available in the library which are untested but further expand the the number of tokenizers by a factor of $44/3$ to 86 million.

All output sequences consist of four token regions representing different features of the maze of which we see an example in Figure 4.

```
<ADJLIST_START> (0,0) <--> (1,0) ; (2,0) <--> (3,0) ; (4,1) <--> (4,0) ; (2,0) <--> (2,1) ;
(1,0) <--> (1,1) ; (3,4) <--> (2,4) ; (4,2) <--> (4,3) ; (0,0) <--> (0,1) ; (0,3) <--> (0,2) ;
(4,4) <--> (3,4) ; (4,3) <--> (4,4) ; (4,1) <--> (4,2) ; (2,1) <--> (2,2) ; (1,4) <--> (0,4) ;
(1,2) <--> (0,2) ; (2,4) <--> (2,3) ; (4,0) <--> (3,0) ; (2,2) <--> (3,2) ; (1,2) <--> (2,2) ;
(1,3) <--> (0,3) ; (3,2) <--> (3,3) ; (0,2) <--> (0,1) ; (3,1) <--> (3,2) ; (1,3) <--> (1,4) ;
<ADJLIST_END>      <ORIGIN_START> (1,3) <ORIGIN_END>      <TARGET_START> (2,3) <TARGET_END>
<PATH_START> (1,3) (0,3) (0,2) (1,2) (2,2) (2,1) (2,0) (3,0) (4,0) (4,1) (4,2) (4,3) (4,4)
(3,4) (2,4) (2,3) <PATH_END>
```

Figure 4: Example text output format with token regions highlighted. Adjacency list : text representation of the lattice graph, Origin : starting coordinate, Target : ending coordinate, Path : maze solution sequence

Each `MazeTokenizerModular` is constructed from a set of several `_TokenizerElement` objects, each of which specifies how different token regions or other elements of the stringification are produced.
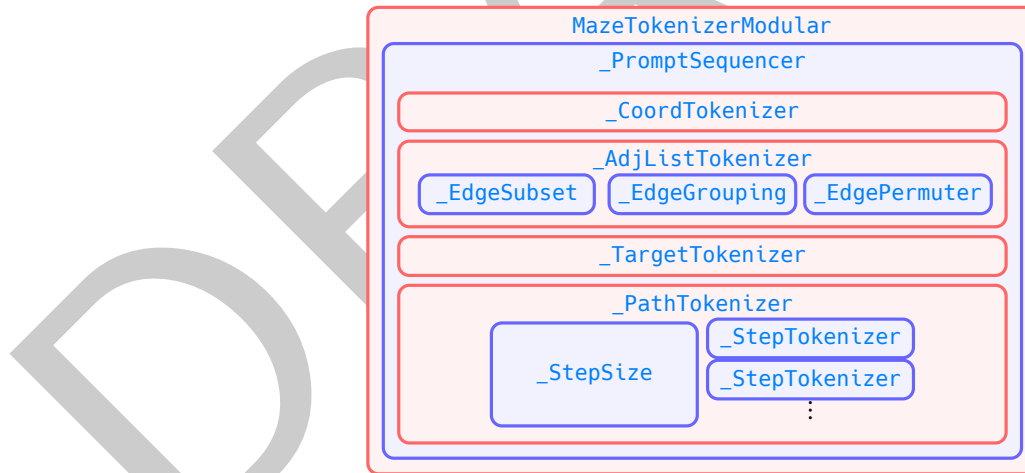
Figure 5: Nested internal structure of `_TokenizerElement` objects inside a typical `MazeTokenizerModular`.

The tokenizer architecture is purposefully designed such that adding and testing a wide variety of new tokenization algorithms is fast and minimizes disturbances to functioning code. This is enabled by the modular architecture and the automatic inclusion of any new tokenizers in integration tests. To create a new tokenizer, developers forking the library may simply create their own `_TokenizerElement` subclass and implement the abstract methods. If the behavior change is sufficiently small, simply adding a parameter to an existing `_TokenizerElement` subclass and updating its implementation will suffice. For small additions, simply adding new cases to existing unit tests will suffice.

The breadth of tokenizers is also easily scaled in the opposite direction. Due to the exponential scaling of parameter combinations, adding a small number of new features can significantly

135 slow certain procedures which rely on constructing all possible tokenizers, such as integration
136 tests. If any existing subclass contains features which aren't needed, a developer tool decorator
137 `@mark_as_unsupported` is provided which can be applied to the unneeded `_TokenizerElement`
138 subclasses to prune those features and compact the available space of tokenizers.

Benchmarks of Generation Speed

140 We provide approximate benchmarks for relative generation time across various algorithms,
141 parameter choices, maze sizes, and dataset sizes in Table 1 and Figure 6. Experiments were
142 performed on a standard GitHub runner without parallelism.

maze_ctor	keyword args	all sizes	small $g \leq 10$	medium $g \in (10, 32]$	large $g > 32$
dfs		28.0	2.8	20.3	131.8
dfs	accessible_cells=20	2.3	2.2	2.4	2.2
dfs	do_forks=False	2.7	2.2	3.1	3.5
dfs	max_tree_depth=0.5	2.5	2.0	2.7	4.0
dfs_percolation	p=0.1	43.9	2.8	33.9	208.0
dfs_percolation	p=0.4	48.7	3.0	36.5	233.5
kruskal		12.8	1.9	10.3	55.8
percolation	p=1.0	50.2	2.6	37.2	242.5
recursive_div		10.2	1.7	8.9	42.1
wilson		676.5	7.8	188.6	3992.6
mean		559.9	13.0	223.5	3146.9
median		11.1	6.5	32.9	302.7

Table 1: Generation times for various algorithms and maze sizes.


Figure 6: Plots of maze generation time. Generation time scales exponentially with maze size for all algorithms (left). Generation time does not depend on the number of mazes being generated, and there is minimal overhead to initializing the generation process for a small dataset (right). Wilson's algorithm is notably less efficient than others and has high variance. Note that for both plots, values are averaged across all parameter sets for that algorithm, and parallelization is disabled.

Success Rate Estimation

144 In order to replicate the exact dataset distribution of (Schwarzschild, Borgnia, Gupta, Bansal, et
145 al., 2021), we allow placing additional constraints in `MazeDatasetConfig.endpoint_kwargs`:

₁₄₆ EndpointKwargsType, such as enforcing that the start or end point be in a "dead end" with
₁₄₇ only one accessible neighbor cell. However, combining this with cyclic mazes (such as those
₁₄₈ generated with percolation), as was required for the work in (Knutson et al., 2024), can lead
₁₄₉ to an absence of valid start and end points. Placing theoretical bounds on this success rate is
₁₅₀ difficult, as it depends on the exact maze generation algorithm and parameters used. However,
₁₅₁ our package provides a way to estimate the success rate of a given configuration using a
₁₅₂ symbolic regression model trained with PySR (Cranmer, 2023). More details on this can be
₁₅₃ found in estimate_dataset_fractions.ipynb.

₁₅₄ Using the estimation simply requires the user to call cfg_new: MazeDatasetConfig =
₁₅₅ cfg.success_fraction_compensate(), providing their initial cfg and then using the returned
₁₅₆ cfg_new in its place.

Success Rate Estimation Algorithm

₁₅₈ The base function learned by symbolic regression is not particularly insightful, and is potentially
₁₅₉ subject to change. It is defined as cfg_success_predict_fn, and takes a 5 dimensional float
₁₆₀ vector created by MazeDatasetConfig._to_ps_array() which represents the 0) percolation
₁₆₁ value 1) grid size 2) endpoint deadend configuration 3) endpoint uniqueness 4) categorical
₁₆₂ generation function index.

₁₆₃ However, the outputs of this function are not directly usable due to minor divergences at the
₁₆₄ endpoints with respect to the percolation probability p. Since we know that maze success is
₁₆₅ either guaranteed or impossible for $p = 0$ and $p = 1$, we define the soft_step function to
₁₆₆ nudge the raw output of the symbolic regression. This function is defined with:

₁₆₇ A shifted sigmoid σ_s, amplitude scaling A, and h function:

$$\sigma_s(x) = (1 + e^{-10^3 \cdot (x-0.5)})^{-1} \qquad A(q, a, w) = w \cdot (1 - |2q - 1|^a)$$

$$h(q, a) = q \cdot (1 - |2q - 1|^a) \cdot (1 - \sigma_s(q)) + (1 - (1 - q) \cdot (1 - |2(1 - q) - 1|^a)) \cdot \sigma_s(q)$$

₁₆₉ We combine these to get the soft_step function, which is identity-like for $p \approx 0.5$, and pushes
₁₇₀ pushes x to extremes otherwise.

$$\text{soft\_step}(x, p, \alpha, w) = h(x, A(p, \alpha, w))$$

₁₇₁ Finally, we define

$$\text{cfg\_success\_predict\_fn}(\mathbf{x}) = \text{soft\_step}(\text{raw\_val}, x_0, 5, 10)$$

₁₇₂ where raw_val is the output of the symbolic regression model. x_0 is the percolation probability,
₁₇₃ while all other parameters from _to_ps_array() only affect raw_val.

Figure 7: An example of both empirical and predicted success rates as a function of the percolation probability p for various maze sizes, percolation with and without depth first search, and `endpoint_kwargs` requiring that both the start and end be in unique dead ends.

Implementation

We refer to our GitHub repository and docs for documentation and up-to-date implementation details.

This package utilizes a simple, efficient representation of mazes. Using an adjacency list to represent mazes would lead to a poor lookup time for whether any given connection exists, while using an adjacency matrix would waste memory by failing to exploit the structure (e.g., only 4 of the diagonals would be filled in). Instead, we describe mazes with the following simple representation: for a d-dimensional lattice with r rows and c columns, we initialize a boolean array $A = \{0, 1\}^{d \times r \times c}$, which we refer to in the code as a `connection_list`. The value at $A[0, i, j]$ determines whether a downward connection exists from node $[i, j]$ to $[i + 1, j]$. Likewise, the value at $A[1, i, j]$ determines whether a rightwards connection to $[i, j + 1]$ exists. Thus, we avoid duplication of data about the existence of connections, at the cost of requiring additional care with indexing when looking for a connection upwards or to the left. Note that this setup allows for a periodic lattice.

To produce solutions to mazes, two points are selected uniformly at random without replacement from the connected component of the maze, and the A^* algorithm (Hart et al., 1968) is applied to find the shortest path between them. The endpoint selection can be affected by `MazeDatasetConfig.endpoint_kwargs: EndpointKwargsType`, and complications caused by this are detailed in subsection .

Parallelization is implemented via the `multiprocessing` module in the Python standard library, and parallel generation can be controlled via keyword arguments to `MazeDataset.from_config()`.

Usage in Research

This package was originally built for the needs of the (Michael I. Ivanitskiy et al., 2023b) project, which aims to investigate spatial planning and world models in autoregressive transformer models trained on mazes (Michael Igorevich Ivanitskiy et al., 2023; Spies et al., 2024). This project has also adapted itself to be useful for work on understanding the mechanisms by which recurrent convolutional and implicit networks (Fung et al., 2022) solve mazes given a rasterized view (Knutson et al., 2024), and for this we match the output format of (Schwarzschild,

Borgnia, Gupta, Bansal, et al., 2021).

This package has also been utilized in work by other groups:

- (Nolte et al., 2024) use maze-dataset to compare the effectiveness of transformers trained with the MLM-$\mathcal{U}$ (Kitouni et al., 2024) multistep prediction objective against standard autoregressive training for multi-step planning on our maze task.

- (Wang et al., 2024) and (Chen et al., 2024) use maze-dataset to study the effectiveness of imperative learning.

Acknowledgements

References

Alance AB. (2019). *Maze generator*. http://www.mazegenerator.net.

Ayaz, H., Allen, S. L., Platek, S. M., & Onaral, B. (2008). Maze suite 1.0: A complete set of tools to prepare, present, and analyze navigational and spatial cognitive neuroscience experiments. *Behavior Research Methods*, *40*, 353–359.

Chen, X., Yang, F., & Wang, C. (2024). iA*: Imperative learning-based A* search for pathfinding. *arXiv Preprint arXiv:2403.15870*.

Cobbe, K., Hesse, C., Hilton, J., & Schulman, J. (2019). Leveraging procedural generation to benchmark reinforcement learning. *arXiv Preprint arXiv:1912.01588*.

Cranmer, M. (2023). Interpretable machine learning for science with PySR and SymbolicRegression. jl. *arXiv Preprint arXiv:2305.01582*.

Duminil-Copin, H. (2017). *Sixty years of percolation* (No. arXiv:1712.04651). arXiv. http://arxiv.org/abs/1712.04651

Ehsan, E. (2022). *Maze*. https://github.com/emadehsan/maze

Fisher, M. E., & Essam, J. W. (2004). Some Cluster Size and Percolation Problems. *Journal of Mathematical Physics*, *2*(4), 609–619. https://doi.org/10.1063/1.1703745

Fung, S. W., Heaton, H., Li, Q., McKenzie, D., Osher, S., & Yin, W. (2022). Jfb: Jacobian-free backpropagation for implicit networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, *36*, 6648–6656.

Guo, C., Barthelet, L., & Morris, R. (2011). *Maze generator and solver*. Wolfram Demonstrations Project, https://demonstrations.wolfram.com/MazeGeneratorAndSolver/.

Harries, L., Lee, S., Rzepecki, J., Hofmann, K., & Devlin, S. (n.d.). MazeExplorer: A Customisable 3D Benchmark for Assessing Generalisation in Reinforcement Learning. *2019*

245   *IEEE Conf. Games CoG*, 1–4.

246 Hart, P. E., Nilsson, N. J., & Raphael, B. (1968). A Formal Basis for the Heuristic Determination
247   of Minimum Cost Paths. *IEEE Transactions on Systems Science and Cybernetics*, *4*(2),
248   100–107. https://doi.org/10.1109/TSSC.1968.300136

249 Ivanitskiy, M. (n.d.). *ZANJ*. https://github.com/mivanit/ZANJ

250 Ivanitskiy, Michael I., Shah, R., Spies, A. F., Räuker, T., Valentine, D., Rager,
251   C., Quirke, L., Corlouer, G., & Mathwin, C. (2023a). *Maze dataset*. https:
252   //github.com/understanding-search/maze-dataset

253 Ivanitskiy, Michael I., Shah, R., Spies, A. F., Räuker, T., Valentine, D., Rager, C., Quirke,
254   L., Corlouer, G., & Mathwin, C. (2023b). *Maze transformer interpretability*. https:
255   //github.com/understanding-search/maze-transformer

256 Ivanitskiy, Michael Igorevich, Spies, A. F., Räuker, T., Corlouer, G., Mathwin, C., Quirke,
257   L., Rager, C., Shah, R., Valentine, D., Behn, C. D., & others. (2023). Structured world
258   representations in maze-solving transformers. *arXiv Preprint arXiv:2312.02566*.

259 Kitouni, O., Nolte, N. S., Williams, A., Rabbat, M., Bouchacourt, D., & Ibrahim, M. (2024).
260   The factorization curse: Which tokens you predict underlie the reversal curse and more.
261   *Advances in Neural Information Processing Systems*, *37*, 112329–112355.

262 Knutson, B., Rabeendran, A. C., Ivanitskiy, M., Pettyjohn, J., Diniz-Behn, C., Fung, S. W.,
263   & McKenzie, D. (2024). On logical extrapolation for mazes with recurrent and implicit
264   networks. *arXiv Preprint arXiv:2410.03020*.

265 Kruskal, J. B. (1956). On the shortest spanning subtree of a graph and the traveling salesman
266   problem. *Proceedings of the American Mathematical Society*, *7*(1), 48–50.

267 Liu, C., & Wu, B. (2023). Evaluating large language models on graphs: Performance insights
268   and comparative analysis. *arXiv Preprint arXiv:2308.11224*.

269 Nag, A. (2020). MDL suite: A language, generator and compiler for describing mazes. *Journal*
270   *of Open Source Software*, *5*(46), 1815.

271 Németh, F. (2019). *Maze-generation-algorithms*. https://github.com/ferenc-nemeth/
272   maze-generation-algorithms

273 Nolte, N., Kitouni, O., Williams, A., Rabbat, M., & Ibrahim, M. (2024). Transformers can
274   navigate mazes with multi-step prediction. *arXiv Preprint arXiv:2412.05117*.

275 Oppenheim, J. (2018). *Maze-generator: Generate a random maze represented as a 2D array*
276   *using depth-first search*. https://github.com/oppenheimj/maze-generator/; GitHub.

277 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
278   Gimelshein, N., Antiga, L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M.,
279   Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., … Chintala, S. (2019). PyTorch: An
280   Imperative Style, High-Performance Deep Learning Library. In H. Wallach, H. Larochelle, A.
281   Beygelzimer, F. d'Alché-Buc, E. Fox, & R. Garnett (Eds.), *Advances in neural information*
282   *processing systems 32* (pp. 8024–8035). Curran Associates, Inc. http://papers.neurips.cc/
283   paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf

284 Pfau, J., Merrill, W., & Bowman, S. R. (2024). Let's think dot by dot: Hidden computation
285   in transformer language models. *arXiv Preprint arXiv:2404.15758*.

286 Räuker, T., Ho, A., Casper, S., & Hadfield-Menell, D. (2023). Toward transparent ai: A survey
287   on interpreting the inner structures of deep neural networks. *2023 IEEE Conference on*
288   *Secure and Trustworthy Machine Learning (SaTML)*, 464–483.

289 Schwarzschild, A., Borgnia, E., Gupta, A., Bansal, A., Emam, Z., Huang, F., Goldblum, M., &
290   Goldstein, T. (2021). *Datasets for Studying Generalization from Easy to Hard Examples*

(No. arXiv:2108.06011). arXiv. https://doi.org/10.48550/arXiv.2108.06011

Schwarzschild, A., Borgnia, E., Gupta, A., Huang, F., Vishkin, U., Goldblum, M., & Goldstein, T. (2021). Can you learn an algorithm? Generalizing from easy to hard problems with recurrent networks. *Advances in Neural Information Processing Systems*, *34*, 6695–6706.

Singla, A. (2023). Evaluating ChatGPT and GPT-4 for visual programming. *arXiv Preprint arXiv:2308.02522*.

Spies, A. F., Edwards, W., Ivanitskiy, M. I., Skapars, A., Räuker, T., Inoue, K., Russo, A., & Shanahan, M. (2024). Transformers use causal world models in maze-solving tasks. *arXiv Preprint arXiv:2412.11867*.

Wang, C., Ji, K., Geng, J., Ren, Z., Fu, T., Yang, F., Guo, Y., He, H., Chen, X., Zhan, Z., & others. (2024). Imperative learning: A self-supervised neural-symbolic learning framework for robot autonomy. *arXiv Preprint arXiv:2406.16087*.

Wilson, D. B. (1996). Generating random spanning trees more quickly than the cover time. *Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing - STOC '96*, 296–303. https://doi.org/10.1145/237814.237880